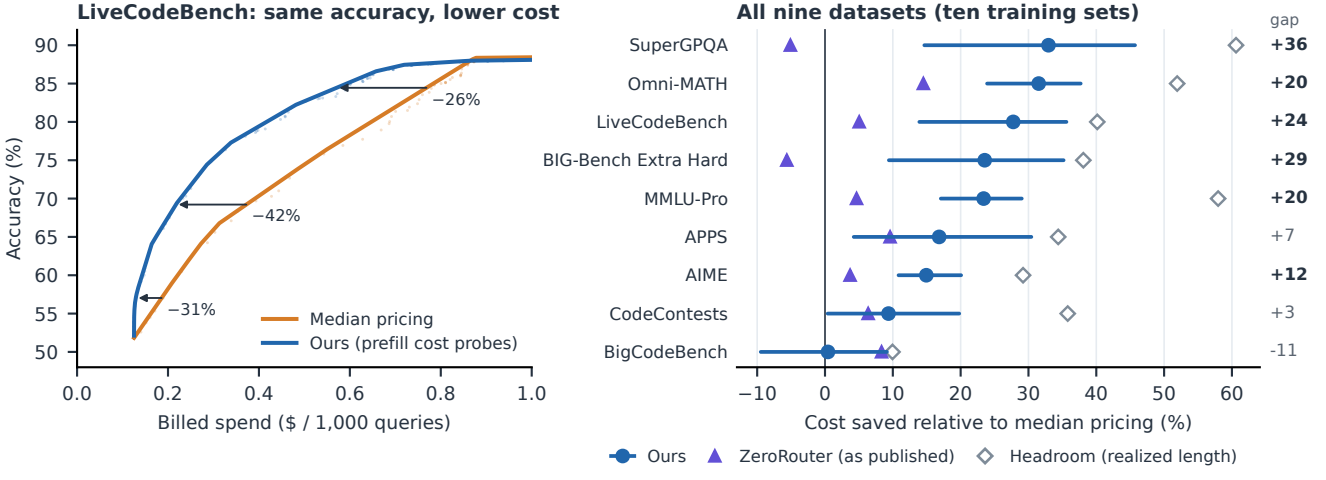
LiveCodeBench: same accuracy, lower cost
Accuracy (%)
90
85
80
75
70
65
60
55
50
0.0
0.2
0.4
0.6
0.8
1.0
Billed spend ($ / 1,000 queries)
−26%
−42%
−31%
Median pricing
Ours (prefill cost probes)
All nine datasets (ten training sets)
gap
SuperGPQA
Omni-MATH
LiveCodeBench
BIG-Bench Extra Hard
MMLU-Pro
APPS
AIME
CodeContests
BigCodeBench
+36
+20
+24
+29
+20
+7
+12
+3
-11
−10
0
10
20
30
40
50
60
Cost saved relative to median pricing (%)
Ours
ZeroRouter (as published)
Headroom (realized length)