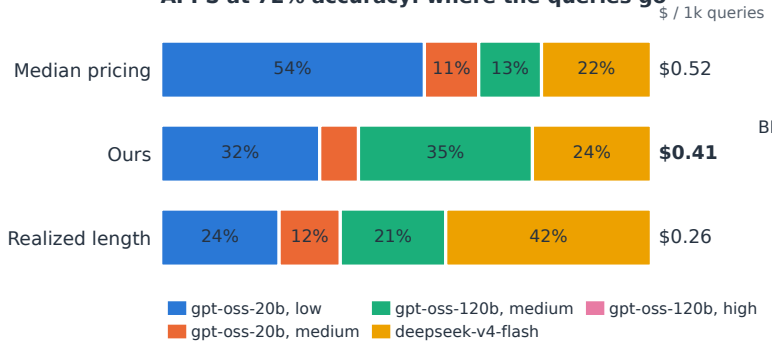
APPS at 72% accuracy: where the queries go
$ / 1k queries
Median pricing
54%
11%
13%
22%
$0.52
Ours
32%
35%
24%
$0.41
Realized length
24%
12%
21%
42%
$0.26
gpt-oss-20b, low
gpt-oss-120b, medium
gpt-oss-120b, high
gpt-oss-20b, medium
deepseek-v4-flash

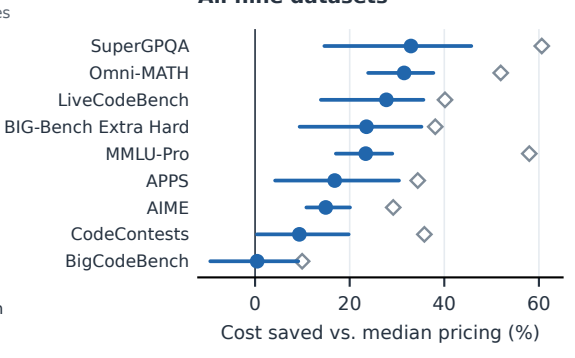
All nine datasets
SuperGPQA
Omni-MATH
LiveCodeBench
BIG-Bench Extra Hard
MMLU-Pro
APPS
AIME
CodeContests
BigCodeBench
0
20
40
60
Cost saved vs. median pricing (%)
Ours, 95% interval
Perfect cost knowledge