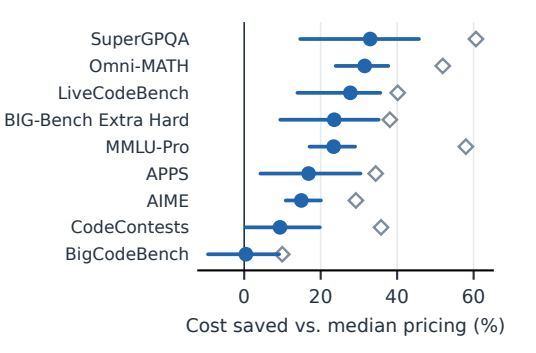
All nine datasets
SuperGPQA
Omni-MATH
LiveCodeBench
BIG-Bench Extra Hard
MMLU-Pro
APPS
AIME
CodeContests
BigCodeBench
0
20
40
60
Cost saved vs. median pricing (%)
Ours, 95% interval
Perfect cost knowledge

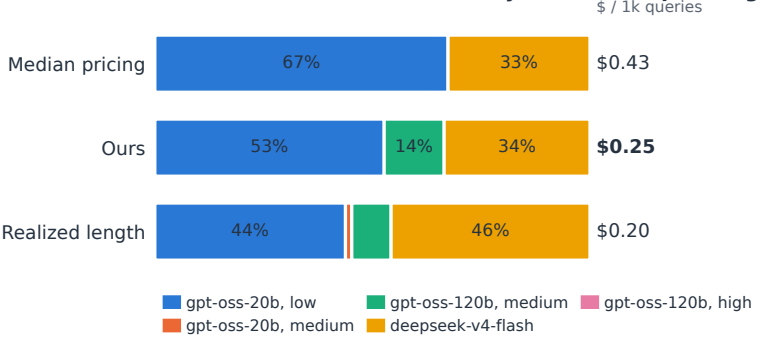
LiveCodeBench at 72% accuracy: where the queries go
$ / 1k queries
Median pricing
67%
33%
$0.43
Ours
53%
14%
34%
$0.25
Realized length
44%
46%
$0.20
gpt-oss-20b, low
gpt-oss-120b, medium
gpt-oss-120b, high
gpt-oss-20b, medium
deepseek-v4-flash