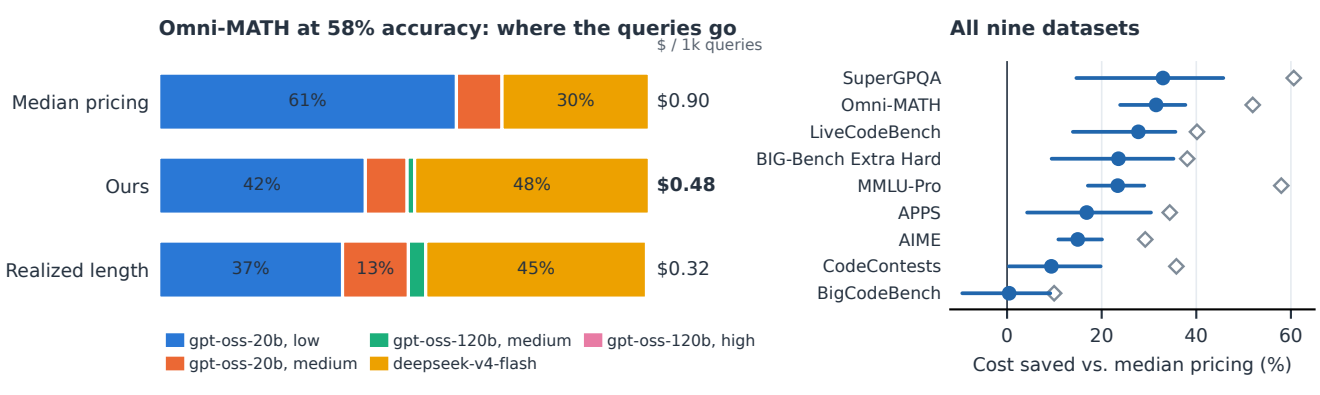
Omni-MATH at 58% accuracy: where the queries go
$ / 1k queries
Median pricing
61%
30%
$0.90
Ours
42%
48%
$0.48
Realized length
37%
13%
45%
$0.32
gpt-oss-20b, low
gpt-oss-120b, medium
gpt-oss-120b, high
gpt-oss-20b, medium
deepseek-v4-flash
All nine datasets
SuperGPQA
Omni-MATH
LiveCodeBench
BIG-Bench Extra Hard
MMLU-Pro
APPS
AIME
CodeContests
BigCodeBench
0
20
40
60
Cost saved vs. median pricing (%)
Ours, 95% interval
Perfect cost knowledge