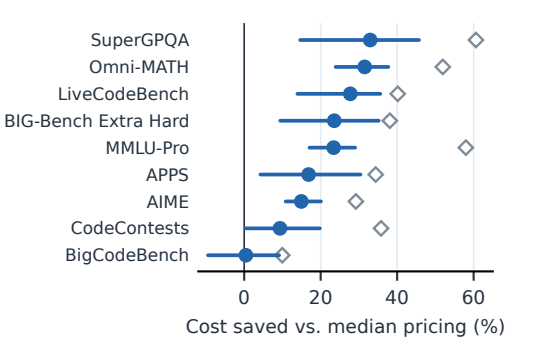
All nine datasets
SuperGPQA
Omni-MATH
LiveCodeBench
BIG-Bench Extra Hard
MMLU-Pro
APPS
AIME
CodeContests
BigCodeBench
0
20
40
60
Cost saved vs. median pricing (%)
Ours, 95% interval
Perfect cost knowledge

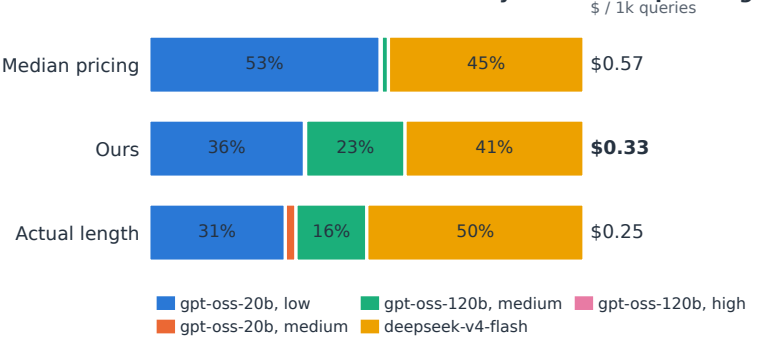
LiveCodeBench at 77% accuracy: where the queries go
$ / 1k queries
Median pricing
53%
45%
$0.57
Ours
36%
23%
41%
$0.33
Actual length
31%
16%
50%
$0.25
gpt-oss-20b, low
gpt-oss-120b, medium
gpt-oss-120b, high
gpt-oss-20b, medium
deepseek-v4-flash