gpt-oss-20b, low
gpt-oss-20b, medium
gpt-oss-120b, medium
deepseek-v4-flash
gpt-oss-120b, high
Share of test queries per route (%), routes ordered by cost per query; right: spend per 1,000 queries
LiveCodeBench
at 72% accuracy; ours saves 42%
Median pricing
Ours
Realized length
67
33
$0.43
53
14
34
$0.25
44
46
$0.20
Omni-MATH
at 58% accuracy; ours saves 46%
61
30
$0.90
42
48
$0.48
37
13
45
$0.32
MMLU-Pro
at 75% accuracy; ours saves 34%
46
52
$0.20
32
53
$0.13
41
55
$0.07
AIME
at 79% accuracy; ours saves 18%
49
51
$0.57
37
63
$0.47
29
64
$0.37
APPS
at 72% accuracy; ours saves 21%
54
11
13
22
$0.52
32
35
24
$0.41
24
12
21
42
$0.26
CodeContests
at 67% accuracy; ours saves 8%
49
19
29
$0.58
34
27
37
$0.54
31
28
31
$0.31
SuperGPQA
at 49% accuracy; ours saves 33%
59
41
$0.20
45
50
$0.13
62
38
$0.07
BIG-Bench Extra Hard
at 45% accuracy; ours saves 37%
68
25
$0.68
45
11
37
$0.43
47
11
34
$0.34
BigCodeBench
at 49% accuracy; ours saves 3%
67
32
$0.07
66
30
$0.07
67
31
$0.06
0
50
100%